

CASE STUDY 2

Ship the trust loop. Cut everything else.

MVP scope, release plan, and the features I'm deliberately not building for GDT's marketplace launch.

THE TEST

Is it in the trust loop?

THE CUT

If no: defer. If yes: manual OK?

THE TIMELINE

Release plans

THE PROBLEM

Five stakeholders want five different MVPs

Someone has to pick the cut line.

What Each Stakeholder Wants

FOUNDER – *Speed*

Ship fast. Anything that moves.

OPERATIONS – *Automation*

Fewer manual steps, less firefighting.

CUSTOMER SUCCESS – *Clarity*

Clear flows so travellers know what's happening.

MARKETING – *Polish*

Polish — the site can't look half-built.

ENGINEERING – *Scope*

Small scope. Capacity is tight.

The Resolution

One Test to Replace all Preferences

The MVP is the smallest unbroken trust loop a traveller can complete end-to-end.

[assumption] 6-person team · 6-week launch window · pre-launch, no live traffic yet.

FRAMEWORK

MoSCoW, filtered by two questions

TWO-QUESTION FILTER (APPLIED TO EVERY BACKLOG ITEM)

Q1 Is this in the trust loop?

The minimum path: discover → evaluate → book → pay → receive → travel → refund path.



Q2 Can it be manual without breaking trust?

Manual = a human does it behind the scenes. Broken trust = the traveller sees the seam.



NO → DEFER

Not in the loop. Ship later.

YES → MANUAL OK

Stub it with ops. Automate later.

YES → MUST AUTOMATE

In the loop AND manual breaks trust.

THE CUT LINE

16 requested areas → 9 Must, 4 Should, 3 Defer

MUST HAVE9	SHOULD HAVE4	DEFER3
In the trust loop. Ship or don't launch. • Vendor onboarding <i>No supply, no marketplace</i> • Listing setup <i>The product itself</i> • Admin approvals <i>Curation is GDT's moat</i> • Discovery + PDP (Product Detail Page) <i>How travellers evaluate</i> • Booking + payment <i>The transaction itself</i> • Cancellation policy (visible) <i>Pre-booking trust signal</i> • Booking confirmation (email) <i>The receipt</i> • Refunds (manual) <i>Money stuck = trust dead</i> • Vendor payouts (manual) <i>No vendors stay without this</i>	Ships in R2, within 4 weeks of launch. • Basic reviews <i>Cold-start problem — wait for volume</i> • Promotions / coupons <i>Growth lever, not launch-critical</i> • Ops analytics <i>Excel works day 1; dashboard can wait</i> • Support ticket visibility <i>WhatsApp sufficient at low volume</i>	Not now. Revisit with data. • In-app messaging <i>WhatsApp is free, ships today</i> • Field Notes deep integration <i>Content enriches discovery but isn't transactional. It is growth lever</i> • Commissions automation <i>Spreadsheet covers <100 bookings/mo</i>

RELEASE PLAN

Each one ships a working trust loop

R0 Walking Skeleton Weeks 1-3 · internal only	R1 Beta Lauch Weeks 4-6 · invite-only launch	R2 Customer launch Weeks 7-10 · public launch
GOAL Can one trusted friend book one curated experience end-to-end? WHAT SHIPS 1 vendor, 3 listings (hand-loaded by ops) - Discovery: category page + listing page only - Booking + payment (Razorpay integration) - Manual: approvals, refunds, payouts, support GATE TO NEXT RELEASE E2E booking works on 3 test devices with zero manual intervention during checkout	GOAL Can 100 strangers book 10+ listings with an ops team of 2? WHAT SHIPS - Self-serve vendor onboarding form - Admin approval queue (internal tool) - Cancellation policy on every PDP - Booking confirmation + reminder emails GATE TO NEXT RELEASE Payment success >95%, refund SLA <48h, support response <4h	GOAL Can we scale to 500 travellers/mo without ops breaking? WHAT SHIPS - Learn from R1 and fix it - Basic reviews (cold-start friendly) - Promotions / Refferals - Ops analytics dashboard - Support ticket tracker GATE TO NEXT RELEASE Manual ops intervention rate drops below 20% of bookings

MANUAL WORKAROUNDS

What we let done with human capacity

Feature	Approach	Rationale
Vendor approval	Ops reviews each onboarding form manually, 1-day SLA	Ensures quality control in early supply; low volume makes manual review feasible and prevents poor experiences entering marketplace
Refunds	Ops triggers refund via Razorpay dashboard on ticket	High familiarity + zero build effort; fast communication without investing in chat infra early
Vendor payouts	Spreadsheet + weekly bank transfer batch	Centralized but simple; allows quick response without building full ticketing system
Traveller-host chat	WhatsApp handoff: ops shares vendor number post-booking	Low scale → manageable manually; helps validate commission logic before automating
Support tickets	Shared Gmail inbox + WhatsApp, tagged by category	Centralized but simple; allows quick response without building full ticketing system
Commissions	Google Sheet: manually reconciled per vendor each week	Low scale → manageable manually; helps validate commission logic before automating
Analytics	Daily Looker Studio or Excel dump from Postgres	Focus on essential metrics only; avoids premature investment in analytics infrastructure

THE PRINCIPLE

Manual is fine when the traveller doesn't see it. Refunds that arrive in 48h feel automated. Approval queues longer than 2 days feel broken. We automate when ops cost, SLA, or volume crosses a measured threshold — not on a hunch.

LAUNCH READINESS

We launch R1 (invite-only) when 6 gates are green. Not before.

TRANSACTION

Payment success rate $\geq 95\%$

Measured across 20 test bookings on Android, iOS Safari, desktop Chrome.

TRANSACTION

Booking confirmation email delivers in $<60s$

For 100% of test bookings. Reminder email fires 24h before experience.

TRUST

Cancellation policy visible on every live PDP

Zero listings published without a cancellation policy attached.

TRUST

Refund SLA: funds back within 48h

Tested end-to-end with ops. Razorpay refund API confirmed working.

OPERATIONS

Ops can handle 10 bookings/day solo

Dry-run: 1 ops person, 10 simulated bookings, resolution time $<4h$ each.

OPERATIONS

Vendor payout runs on dry-run batch

Test with 3 vendors. Bank transfers reconcile to Google Sheet cleanly.

All 6 gates green $\rightarrow$ invite-only launch. Public launch waits for 2 weeks of stable R1 metrics.

RISKS, TRADE-OFFS, METRICS

What I'm trading, what could break it, and how we'll know.

SUCCESS METRICS

Tracked weekly from R1

- **TRUST**
Refund rate <5%; complaint rate <3 per 100 bookings
- **TRANSACTION**
Payment success ≥95%; E2E booking completion ≥60%
- **OPS**
Manual intervention rate, trending down week-on-week
- **SUPPLY**
Active listings, week-on-week growth

TOP RISKS

Likely, high-impact

- **OPS OVERLOAD**
R1 volume exceeds 2-person capacity. Mitigate: pause marketing, don't onboard new vendors until caught up.
- **PAYMENT FAILURE CLUSTER**
Razorpay edge case on specific banks. Mitigate: UPI + card fallback, tested on top 5 Indian banks.
- **VENDOR QUALITY DRIFT**
Curation slips under onboarding pressure. Mitigate: hard cap of 5 new vendors/week in R1.
- **REFUND PIPELINE BREAKS**
Manual refunds take >48h. Mitigate: daily refund queue review, escalation to founder if SLA slips.

WHAT I'M TRADING

And why it's defensible

- **NO POLISH**
Marketing loses. Accepted: polish without a working loop is lipstick.
- **NO AUTOMATION**
Ops hates it at R1. Accepted: we measure thresholds and automate on signal, not preference.
- **NO REVIEWS AT LAUNCH**
Cold-start. Seeding them is dishonest. Accepted: ship curation as the trust proxy.
- **NO FIELD NOTES INTEGRATION**
CS1's Fit Score is Phase 2. Launch first, moat second.

DECISION: Ship R0 walking skeleton in 3 weeks; invite-only R1 by week 6. **VALIDATE:** 6 launch gates, not a date. **NEXT:** Score full backlog against the two-question filter with eng lead this week.